

Chapter 11 System Monitoring - Notes

11.3 Learning Objectives:

- Understand the concept of inventory and gain familiarity with available system monitoring tools.
- Understand where the system stores log files and examine the most important ones.
- Use the `/proc` and `/sys` pseudo-file systems.
- Use `sar` to gather system activity and performance data and create reports that are readable by humans.

11.4 Available Monitoring Tools

Linux distributions come with many standard performance/profiling tools already installed. Many familiar from other UNIX-like operating systems, while some developed specifically for Linux.

Most tools make use of mounted **pseudo-file systems**, especially `/proc` and `/sys`, both of which have already been discussed while examining file systems/kernel configuration. Will look at both.

Also a number of graphical system monitors that hide many details, but will only consider command line tools in course.

Before considering main utilities in detail, can see summary on next few sections, broken down by type: note: some utilities have overlapping domains of coverage. Will revisit tables in following chapters that focus on specific topics.

Summary of main process monitoring utility tools:

Process and Load Monitoring Utilities

Utility	Purpose	Package
top	Process activity, dynamically updated	procps
uptime	How long system is running and average load	procps
ps	Detailed information about processes	procps
pstree	Tree of processes and their connections	psmisc (or pstree)
mpstat	Multiple processor usage	sysstat
iostat	CPU utilization and I/O statistics	sysstat
sar	Display and collect information about system activity	sysstat
numstat	Information about NUMA (Non-Uniform Memory Architecture)	numactl
strace	Information about all system calls a process makes	strace

Memory Monitoring Utilities

Utility	Purpose	Package
free	Brief summary of memory usage	procps

vmstat	Detailed virtual memory statistics and block I/O, dynamically updated	procps
pmap	Process memory map	procps

I/O Monitoring Utilities

Utility	Purpose	Package
iostat	CPU utilization and I/O statistics	sysstat
sar	Display and collect information about system activity	sysstat
vmstat	Detailed virtual memory statistics and block I/O, dynamically updated	procps

Network Monitoring Utilities

Utility	Purpose	Package
netstat	Detailed networking statistics	netstat
iptraf	Gather information on network interfaces	iptraf
tcpdump	Detailed analysis of network packets and traffic	tcpdump
wireshark	Detailed network traffic analysis	wireshark

11.5 System Log Files

System log files -> essential for monitoring/troubleshooting. In Linux, messages appear in various files under `/var/log`. Exact names vary with Linux distribution.

Ultimate control of how messages dealt with -> controlled by **syslogd** (usually **rsyslogd** on modern systems) daemon, common to many UNIX-like operating systems. Newer **systemd**-based systems can use **journalctl** instead, but usually retain **syslogd** and cooperate with it.

Important messages sent not only to logging files, but also to system console window. If not running **X**, or are at virtual terminal, will see them directly there as well. In addition, messages will be copied to `/var/log/messages` (or to `/var/log/syslog` on Ubuntu), but if running **X**, have to take some steps to view them.

Can view new messages continuously as new lines appear with:

```
$ sudo tail -f /var/log/messages (or /var/log/syslog)
```

or

```
$ dmesg -w
```

which shows only kernel-related messages.

11.6 Important Log Files in /var/log

Besides looking at log messages in terminal window, can see them using graphical interfaces.

On GNOME desktop, can also access messages by clicking on `System -> Administration -> System Log` Or `Applications -> System Tools -> Log File` Viewer in your Desktop menus, and other desktops have similar links you can locate.

Some important log files found under `/var/log` :

File	Purpose
<code>boot.log</code>	System boot messages
<code>dmesg</code>	Kernel messages saved after boot. To see current contents of kernel message buffer, type dmesg .
<code>messages</code> Or <code>syslog</code>	All important system messages
<code>secure</code>	Security related messages

In order to keep log files from growing without bound, **logrotate** program run periodically, keeps four previous copies (by default) of log files (optionally compressed). Controlled by `/etc/logrotate.conf`.

11.7 The `/proc` and `/sys` Pseudo-file systems

`/proc` and `/sys` pseudo-file systems contain lot of information about system. Many entries in these directory trees are writable, can be used to change system behavior. Most cases, requires **root** user.

Pseudo-file systems because they totally exist in memory. If look at disk partition when system not running, there will be only empty directory which is used as mount point.

Information displayed is gathered only when looked at. No constant/periodic polling to update entries.

11.8 `/proc` Basics

`/proc` pseudo-file system: long history. Has roots in other UNIX operating system variants. Originally developed to display information about **processes** on system, each of which has own subdirectory in `/proc` with all important process characteristics available.

Over time, grew to contain lot of information about system properties, eg. interrupts, memory, networking, etc. in somewhat anarchistic way. Still extensively used, will often refer to it.

11.9 A survey of `/proc`

What resides in `/proc` pseudo-file system:

```
student@ubuntu: ~  
student@ubuntu:~$ ls -F /proc  
1/      13/      200/     2288/    250/     33/      40/      5383/    99/      misc  
10/     137/     201/     229/     251/     3315/    4047/    54/      990/     modules  
100/     14/      202/     23/      2529/    3336/    41/      5410/    ACPI/     mounts@  
1005/    1457/    203/     230/     2572/    3353/    4156/    5436/    buddyinfo mpt/  
1007/    1465/    204/     231/     26/      3366/    4169/    5438/    bus/      mtrr  
1008/    1478/    205/     232/     27/      3394/    418/     581/     cgroups   net@  
1009/    15/      206/     233/     274/     34/      419/     583/     cmdline   pagetypeinfo  
101/     16/      207/     234/     2756/    3419/    42/      6/       consoles  partitions  
1010/    1649/    208/     235/     276/     3430/    420/     601/     cpuinfo    sched_debug  
1011/    1780/    209/     236/     28/      3439/    422/      7/       crypto     schedstat  
1012/    1782/    21/      2368/    280/     3441/    43/      715/     devices    scsi/  
1013/    1798/    210/     237/     281/     35/      44/      717/     diskstats  self@  
1014/    18/      211/     2373/    2979/    3512/    45/      718/     dma         slabinfo  
102/     1822/    212/     238/     298/     3513/    4595/    721/     driver/     softirqs  
1028/    188/     213/     2385/    2989/    3514/    4596/    723/     execdomains stat  
103/     189/     214/     239/     30/      3515/    4599/    728/     fb          swaps  
104/     19/      2142/    24/      300/     3532/    46/      731/     filesystems sys/  
105/     190/     215/     240/     3081/    3534/    4601/    733/     fs/         sysrq-trigger  
106/     191/     216/     241/     31/      3569/    47/      735/     interrupts sysvipc/  
107/     1916/    217/     242/     32/      3581/    473/     743/     iomem       thread-self@  
108/     192/     218/     243/     3228/    3589/    478/     746/     ioports     timer_list  
1083/    1922/    219/     2436/    3243/    36/      4792/    750/     irq/         timer_stats  
1085/    193/     2196/    244/     3244/    3600/    483/     753/     kallsyms    tty/  
109/     1930/    2197/    2443/    3245/    3613/    485/     755/     kcore        uptime  
11/      194/     22/      2449/    3246/    3621/    486/     781/     keys         version  
110/     1940/    220/     245/     3247/    3655/    492/     798/     key-users    version_signature  
1121/    195/     221/     2456/    3251/    37/      50/      8/      kmsg         vmallocinfo  
1134/    196/     222/     246/     3255/    370/     502/     828/     kpagecgroup vmstat  
1140/    197/     223/     2461/    3257/    38/      51/      9/      kpagecount   zoneinfo  
1141/    198/     224/     247/     3264/    39/      52/      96/     kpageflags  
1147/    1986/    225/     248/     3268/    3910/    53/      98/     loadavg  
116/     199/     226/     2488/    3272/    393/     5310/    981/     locks  
1180/    2/       227/     249/     3274/    399/     5332/    983/     mdstat  
12/     20/      228/     25/      3294/    4/       5358/    988/     meminfo
```

First, see there is subdirectory for each process on system, whether sleeping, running, scheduled out. Looking at random one:

```
student@ubuntu: ~  
student@ubuntu:~$ ls -F /proc/3589  
attr/      coredump_filter  gid_map      mountinfo     oom_score      schedstat      status  
autogroup  cpuset           io           mounts         oom_score_adj  sessionid      syscall  
auxv       cwd@             limits       mountstats    pagemap        setgroups      task/  
cgroup     environ          loginuid     net/           personality     smaps          timers  
clear_refs exe@             map_files/   ns/            projid_map      stack          timerslack_ns  
cmdline    fd/             maps         numa_maps     root@           stat            uid_map  
comm       fdinfo/         mem          oom_adj        sched           statm           wchan
```

Directory full of information about status of process and resources it is using. For example:

```

student@ubuntu: ~
student@ubuntu:~$ cat /proc/3589/status
Name:   bash
Umask:  0022
State:  S (sleeping)
Tgid:   3589
Ngid:   0
Pid:    3589
PPid:   3581
TracerPid: 0
Uid:    1000    1000    1000    1000
Gid:    1000    1000    1000    1000
....
Cpus_allowed:  ffffffff,ffffffff,ffffffff,ffffffff
Cpus_allowed_list:  0-127
Mems_allowed:  00000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00
000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00
000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00
000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00000000,00
0000001
Mems_allowed_list:  0
voluntary_ctxt_switches: 1328
nonvoluntary_ctxt_switches: 100

student@ubuntu:~$

```

Other entries give system-wide information. Eg. can see **interrupt** statistics in below screenshot. For each interrupt, see what type it is, how many times it has been handled on each CPU, which devices registered to respond to it. Also get global statistics.

```

File Edit View Search Terminal Help
x7:/home/coop>cat /proc/interrupts
CPU0      CPU1      CPU2      CPU3
0:         88          0          0          0  IR-I/O-APIC  2-edge    timer
1:        566          1       2153          0  IR-I/O-APIC  1-edge    i8042
8:          1          0          0          0  IR-I/O-APIC  8-edge    rtc0
9:       17990        11        552        27  IR-I/O-APIC  9-fasteoi  acpi
12:      64336        21     186157        20  IR-I/O-APIC 12-edge    i8042
16:          0          0          0          0  IR-I/O-APIC 16-fasteoi i801_smbus
120:         0          0          0          0  DMAR-MSI    0-edge    dmar0
121:         0          0          0          0  DMAR-MSI    1-edge    dmar1
122:     217295       1607        4039     59571  IR-PCI-MSI 376832-edge ahci[0000:00:17.0]
123:          3          0          46          0  IR-PCI-MSI 514048-edge snd_hda_intel:card0
124:     95360        74     49535        192  IR-PCI-MSI 327680-edge xhci_hcd
125:    591286          3     272983          2  IR-PCI-MSI 32768-edge i915
126:         84         16         149         20  IR-PCI-MSI 520192-edge enp0s31f6
127:    225390        29        383         46  IR-PCI-MSI 2097152-edge iwlwifi
NMI:         24        120        130        119  Non-maskable interrupts
LOC:    2255506     2201465     2360863     2239138  Local timer interrupts
SPU:          0          0          0          0  Spurious interrupts
PMI:         24        120        130        119  Performance monitoring interrupts
IWI:          0          0          3          0  IRQ work interrupts
RTR:         24          3          0          0  APIC ICR read retries
RES:    185530     146421     95420     45924  Rescheduling interrupts
CAL:     76456     74989     78143     76063  Function call interrupts
TLB:     75401     73603     76833     75025  TLB shootdowns
ERR:          0
MIS:          0
PIN:          0          0          0          0  Posted-interrupt notification event
PIW:          0          0          0          0  Posted-interrupt wakeup event
x7:/home/coop>

```

11.10 /proc/sys

Most tunable system parameters can be found in subdirectory tree rooted at `/proc/sys` :

```
student@linux-mint ~  
File Edit View Search Terminal Help  
student@linux-mint ~ $ ls -lF /proc/sys  
total 0  
dr-xr-xr-x 1 root root 0 May 31 10:19 abi/  
dr-xr-xr-x 1 root root 0 May 31 10:19 debug/  
dr-xr-xr-x 1 root root 0 May 31 10:19 dev/  
dr-xr-xr-x 1 root root 0 May 31 10:18 fs/  
dr-xr-xr-x 1 root root 0 May 31 10:18 kernel/  
dr-xr-xr-x 1 root root 0 May 31 10:18 net/  
dr-xr-xr-x 1 root root 0 May 31 10:19 sunrpc/  
dr-xr-xr-x 1 root root 0 May 31 10:18 vm/  
student@linux-mint ~ $
```

Each subdirectory contains information + knobs that can be tuned (with care):

- `abi/` : Contains files with application binary information; rarely used
- `debug/` : Debugging parameters; for now, just some control of exception reporting
- `dev/` : Device parameters, including subdirectories for **cdrom**, **scsi**, **raid**, **parport**
- `fs/` : Filesystem parameters, including quota, files handles used, and maximums, inode and directory information, etc.
- `kernel/` : Kernel parameters. Many important entries here.
- `net/` : Network parameters. Subdirectories for **ipv4**, **netfilter**, etc.
- `vm/` : Virtual memory parameters. Many important entries here.

Viewing/changing parameters can be done with simple commands. Eg. maximum number of threads allowed on system seen by looking at:

```
$ ls -l /proc/sys/kernel/threads-max  
$ cat /proc/sys/kernel/threads-max  
129498
```

Can then modify value, verify change was effected:

```
$ sudo bash -c 'echo 100000 > /proc/sys/kernel/threads-max'  
$ cat /proc/sys/kernel/threads-max  
100000
```

Remember from discussion of **sysctl**, same effect accomplished by:

```
$ sudo sysctl kernel.threads-max=100000
```

Viewing value can be done as normal user, changing requires superuser privilege.

11.11 /sys Basics

`/sys` pseudo-filesystem: integral part of **Unified Device Model**. Conceptually, based on **device tree**, one can walk through it and see buses, devices, etc. Also now contains information which may or may not be strictly related to devices, such as kernel modules.

Has more tightly defined structure than `/proc`. Most entries contain only one line of text (although there are exceptions) unlike precursor which has many multi-line entries whose exact contents may change between kernel versions. Thus, interface hopefully more stable.

There are system properties which have display entries in both `/proc` and `/sys`. For compatibility with widely used system utilities, older forms only gradually being whittled down.

11.12 A Survey of `/sys`

Support for **sysfs** virtual filesystem built into all modern kernels, should be mounted under `/sys`. However, unified device model does not require mounting **sysfs** in order to function.

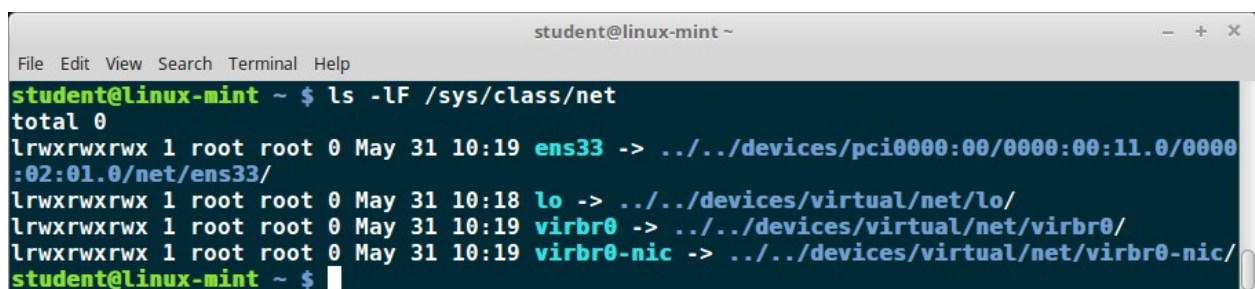
Taking look at 3.18 kernel (warning; exact layout of this filesystem tends to mutate). Top level directory command yields:

```
$ ls -F /sys
block/ bus/ class/ dev/ devices/ firmware/ fs/ kernel/ module/ power/
```

which displays basic device hierarchy. Device model **sysfs** implementation also includes information not strictly related to hardware.

Network devices examined with:

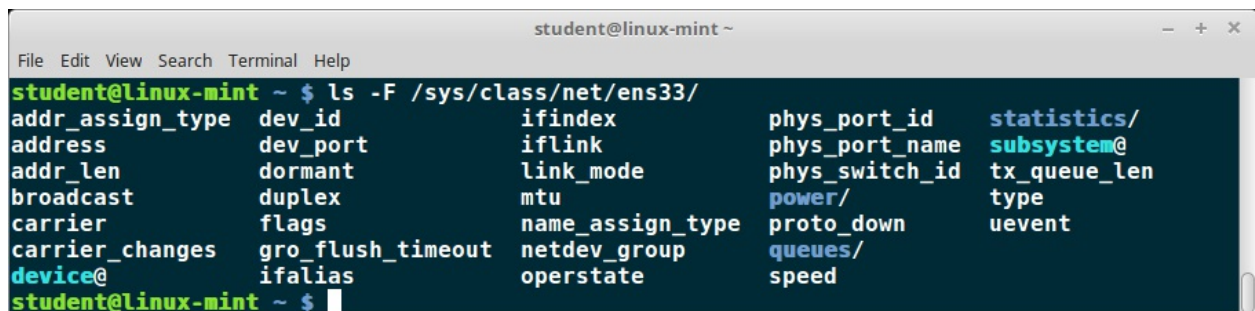
```
$ ls -lF /sys/class/net
```



```
student@linux-mint ~
File Edit View Search Terminal Help
student@linux-mint ~ $ ls -lF /sys/class/net
total 0
lrwxrwxrwx 1 root root 0 May 31 10:19 ens33 -> ../../devices/pci0000:00/0000:00:11.0/0000:02:01.0/net/ens33/
lrwxrwxrwx 1 root root 0 May 31 10:18 lo -> ../../devices/virtual/net/lo/
lrwxrwxrwx 1 root root 0 May 31 10:19 virbr0 -> ../../devices/virtual/net/virbr0/
lrwxrwxrwx 1 root root 0 May 31 10:19 virbr0-nic -> ../../devices/virtual/net/virbr0-nic/
student@linux-mint ~ $
```

Below, can see what looking at Ethernet card gives.

Intention with **sysfs** to have one text value per line, although not expected to be rigorously enforced.



```
student@linux-mint ~
File Edit View Search Terminal Help
student@linux-mint ~ $ ls -lF /sys/class/net/ens33/
addr_assign_type dev_id ifindex phys_port_id statistics/
address dev_port iflink phys_port_name subsystem@
addr_len dormant link_mode phys_switch_id tx_queue_len
broadcast duplex mtu power/
carrier flags name_assign_type proto_down type
carrier_changes gro_flush_timeout netdev_group queues/
device@ ifalias operstate speed
student@linux-mint ~ $
```

Underlying device and driver for first network interface can be traced through `device` and (to be seen shortly) `driver` symbolic links. Below shows what can be seen when looking at directory corresponding to first Ethernet card.

To see full spectrum of information available with **sysfs**, will just have to examine.

```

student@linux-mint ~
File Edit View Search Terminal Help
student@linux-mint ~ $ ls -F /sys/class/net/ens33/device/
acpi_index          dma_mask_bits      local_cpulist      remove            rom
broken_parity_status driver@            local_cpus         rescan            subsystem@
class              driver_override    modalias           reset             subsystem_device
config             enable            msi_bus           resource          subsystem_vendor
consistent_dma_mask_bits firmware_node@     net/              resource0         uevent
d3cold_allowed     irq              numa_node         resource2         vendor
device            label            power/           resource4
student@linux-mint ~ $

```

11.13 sar

sar: Systems Activity Reporter. All-purpose tool for gathering system activity + performance data, creating reports readable by humans.

On Linux systems, backend to **sar** is **sadc** (system activity data collector) which actually accumulates statistics. Stores information in `/var/log/sa` directory, with daily frequency by default, but which can be adjusted. Data collection can be started from command line, regular periodic collection usually started as **cron** job stored in `/etc/cron.d/sysstat`.

sar then reads in this data (either from default locations or by use of file specified with `-f` option), then produces report.

sar invoked via:

```
$ sar [ options ] [ interval ] [ count ]
```

where report repeated after interval seconds a total of count times (defaults to 1). With no options, gives report on CPU usage.

```

student@ubuntu: ~
student@ubuntu:~$ sudo sar 3 3
Linux 4.10.0-20-generic (ubuntu)      06/02/2017      _x86_64_      (4 CPU)

10:13:31 AM    CPU    %user   %nice   %system   %iowait   %steal   %idle
10:13:34 AM    all     77.83    7.08    14.92     0.08     0.00     0.08
10:13:37 AM    all     74.50   12.08    13.42     0.00     0.00     0.00
10:13:40 AM    all     70.56   14.76    14.68     0.00     0.00     0.00
Average:      all     74.30   11.31    14.34     0.03     0.00     0.03
student@ubuntu:~$

```

List of major **sar** options, or modes, each one of which has its own sub-options:

sar Options

Option	Meaning
-A	Almost all Information
-b	I/O and transfer rate statistics (similar to iostat)
-B	Paging statistics including page faults
-x	Block device activity (similar to iostat -x)
-n	Network statistics
-P	Per CPU statistics (as in <code>sar -P ALL 3</code>)

-q	Queue lengths (run queue, processes, and threads)
-r	Swap and memory utilization statistics
-R	Memory statistics
-u	CPU utilization (default)
-v	Statistics about inodes and files and files handles
-w	Context switching statistics
-W	Swapping statistics, pages in and out per second
-f	Extract information from specified file, created by the -o option
-o	Save readings in the file specified, to be read in later with -f option

For example, below can take look at getting paging statistics, and then I/O and transfer rate statistics.

ksar program -> **java**-based utility for generating nice graphs for **sar** data. Can be downloaded from <https://sourceforge.net/projects/ksar/>.

```

student@ubuntu: ~
student@ubuntu:~$ # GETTING PAGING STATISTICS
student@ubuntu:~$
student@ubuntu:~$ sar -B 3 3
Linux 4.10.0-20-generic (ubuntu)      06/02/2017      _x86_64_      (4 CPU)

10:21:44 AM   pgpgin/s   pgpgout/s   fault/s   majflt/s   pgfree/s   pgscank/s   pgscand/s   pgsteal/s   %vmeff
10:21:47 AM      232.00      2117.33  118496.00      0.00  119913.67      0.00      0.00      0.00      0.00
10:21:50 AM      122.67      2853.33  112109.00      0.00  114345.67      0.00      0.00      0.00      0.00
10:21:53 AM      346.67      7170.67  131357.00      0.00  145063.33      0.00      0.00      0.00      0.00
Average:        233.78      4047.11  120654.00      0.00  126440.89      0.00      0.00      0.00      0.00
student@ubuntu:~$
student@ubuntu:~$ # GETTING I/O AND TRANSFER RATE STATISTICS
student@ubuntu:~$
student@ubuntu:~$ sar -b 3 3
Linux 4.10.0-20-generic (ubuntu)      06/02/2017      _x86_64_      (4 CPU)

10:22:01 AM      tps       rtps       wtps   bread/s   bwrtn/s
10:22:04 AM      85.00       22.00       63.00   538.67   9466.67
10:22:07 AM      22.67       14.67        8.00   384.00   4365.33
10:22:10 AM      61.00       10.00       51.00   328.00  65261.33
Average:         56.22       15.56       40.67   416.89  26364.44
student@ubuntu:~$

```

##

[Back to top](#)

[Previous Chapter](#) - [Table of Contents](#) - [Next Chapter](#)